A SELF-ADAPTIVE FITNESS CLUB APPLICATION

Team1: Caules Ge, Ethan Li

What is it?

- A fictional web system for a chain of fitness clubs.
- It supports two user roles: administrators and members.
- Administrators can perform CRUD operations for studios and their affiliated amenities.
- Members can sign up, log in, view available studios, and schedule courses.

Features include an interactive map for studios, and the list is sorted by distance...

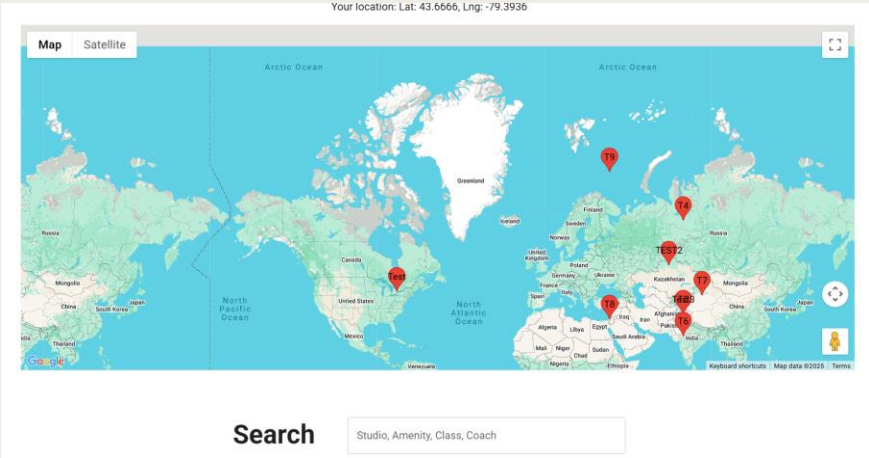

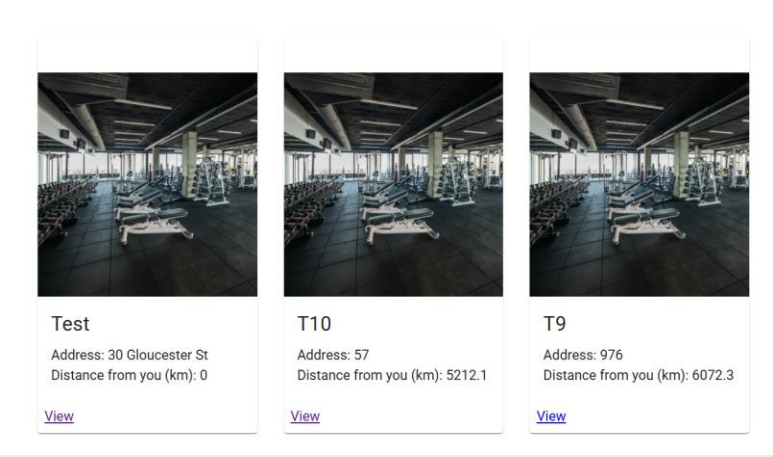

Problem

The Canada Fitness Club system experiences significant traffic surges during promotions.

High traffic causes **slow responses** and **increased errors**, while low traffic results in **wasted resources** due to extra replicas.

Without self-adaptation, the system cannot remain efficient or provide a stable user experience.

Object

3 Layers of Adaptation

- **Client**: simplifies the UI under poor network conditions
- **Server**: scales horizontally based on workload
- **Application**: switches between standard and degraded modes when the workload is extremely high

Goals

Improve response time, reduce errors, and avoid resource waste during low-traffic periods

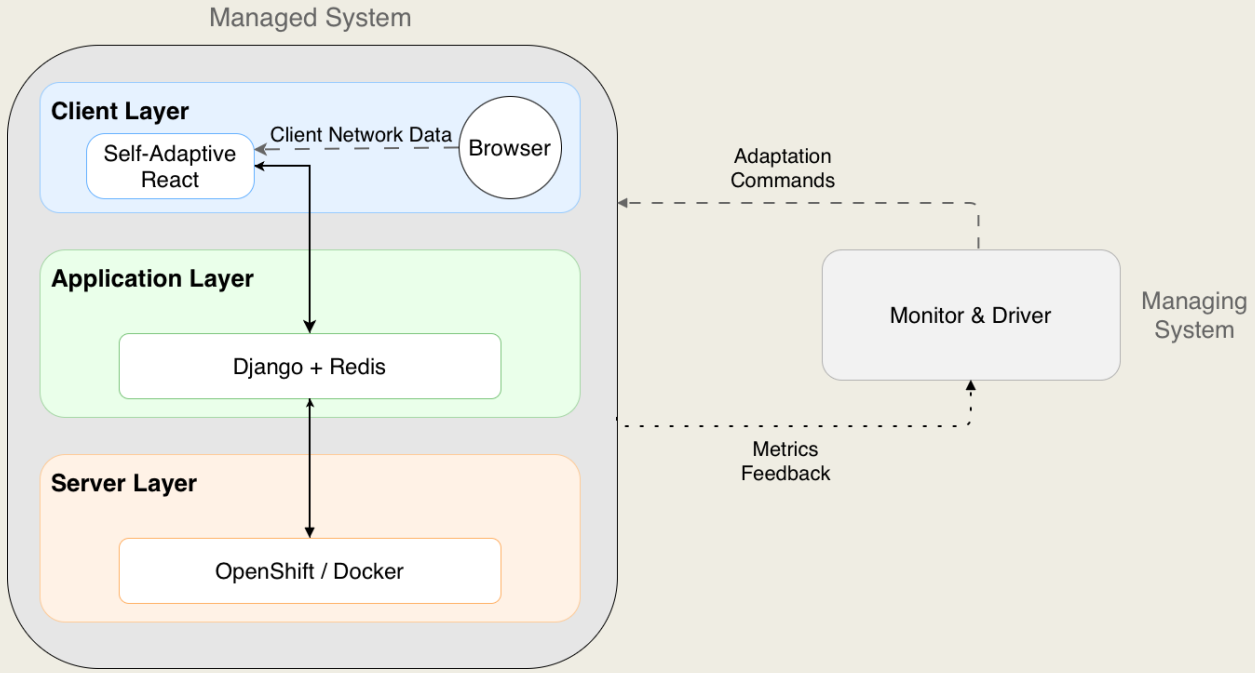

Tools

- Sysdig: collects CPU, memory, error rate, and p95 latency metrics.
- Redis: stores degraded-mode flags and cached responses of each endpoint.
- React: monitors the average FPS, network quality, and kbps download speed of the client. If client

Client-level adaptation:

When the average FPS of the last 5 seconds drops or the network is weak (eg, 3G), the UI switches to a lightweight mode with reduced pagination size, hides the map, and uses lower resolution images

System and Application Level

- When the system scales up and all backend pods are running at maximum capacity, we trigger **Application Adaptation** and switch into **Degraded Mode**

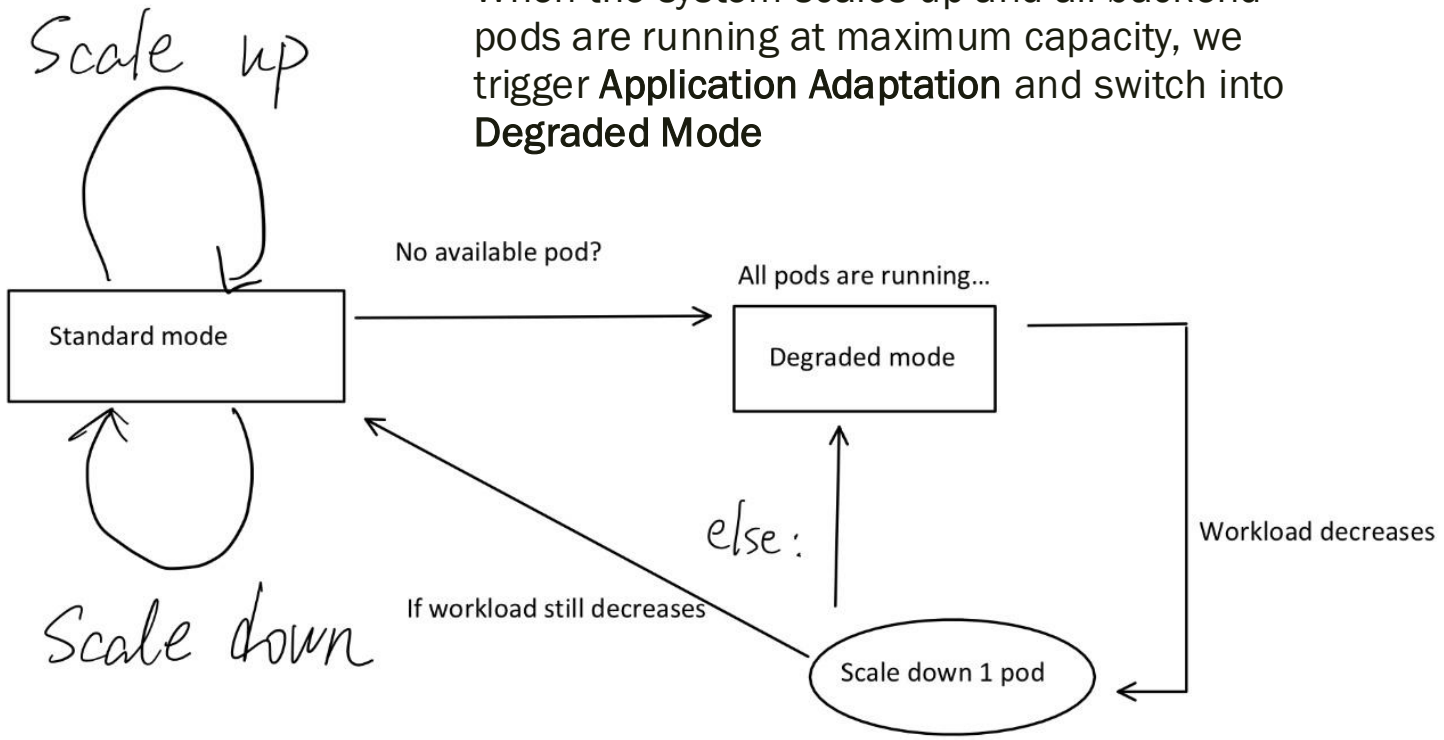

- Reason: avoiding returning to standard mode while all pods are still overloaded, which would harm performance

Degraded mode:

Disabling expensive user-specific features, such as distance-based sorting, allowing Redis caching for pages that normally cannot be cached

For other pages that have been cached, it increases the cache TTL.

Reducing the pagination size from 9 to 6 (fewer DB queries per page)

Result: System and Application Level

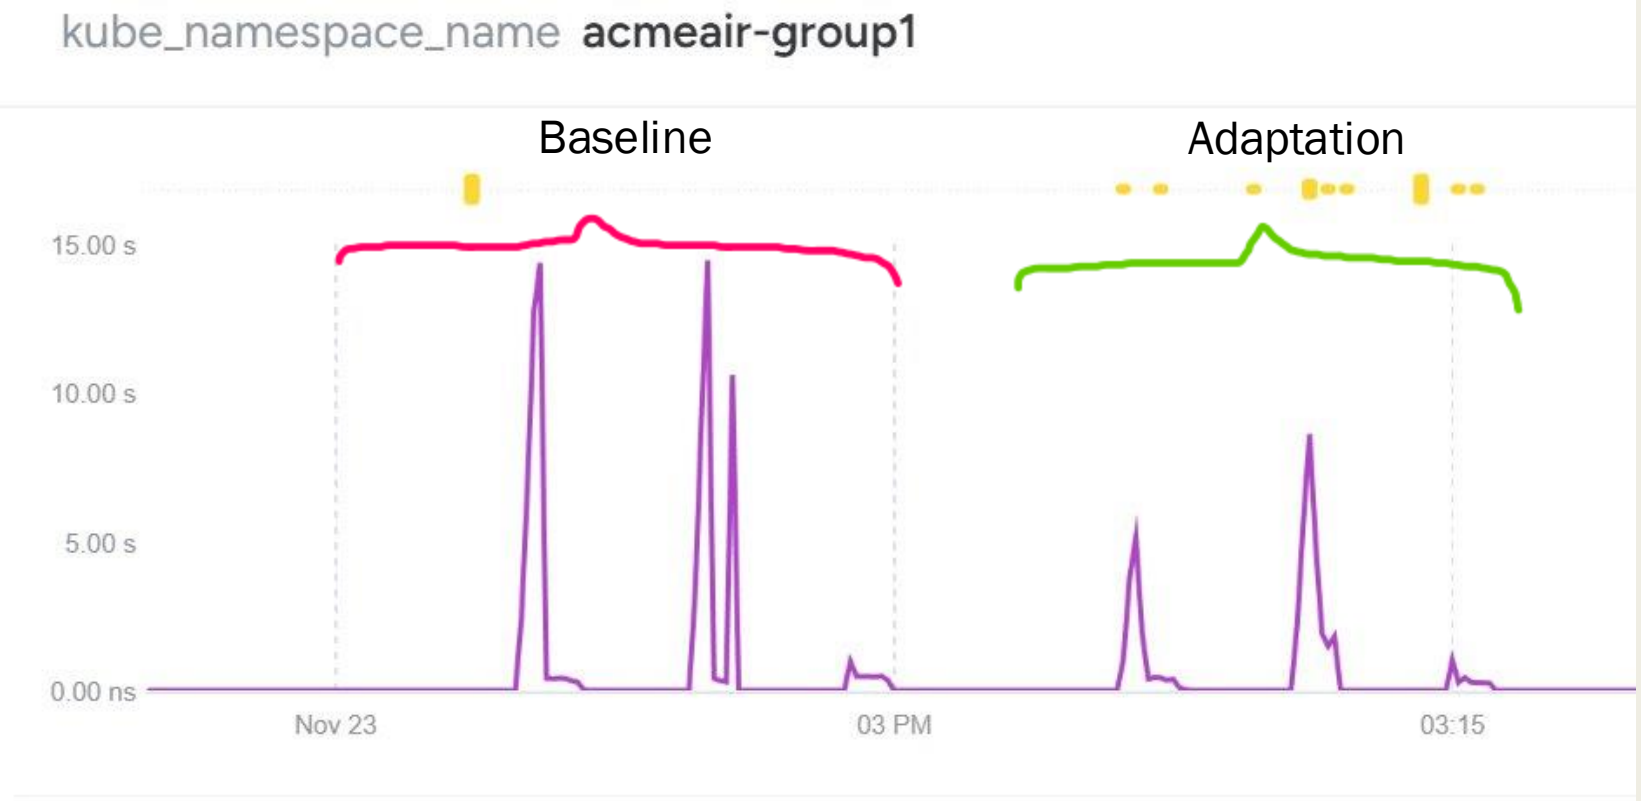

Average response time under three workload levels, from high to low